

## QMF2\_3\_Modelling\_Regression

### Estimating a Linear Regression Model with OLS

This activity introduces ordinary least squares (OLS) linear regression in MATLAB using a tiny, didactic house-pricing dataset. We'll keep the workflow practical.

**(A short theoretical background is provided at the end of the Live Script.)**

Workflow (the basic one):

- Load a small table of house prices and basic features
- Fit a linear model with `fitlm`
- Inspect fit quality with the built-in diagnostic plot
- Make a point prediction for a new house profile using `predict`

#### Dataset

HousePrices is a table file (XLSX) with a numeric response variable Price and a couple of predictors (Size and a simple indicator feature).

This example shows the typical workflow for linear regression analysis using `fitlm`. The general workflow should include preparing a data set, fitting a linear regression model, evaluating and improving the fitted model, and predicting response values for new predictor data.

In the present example we perform only the very basic steps

```
% Import data from spreadsheet
HousePrices2 = readtable("C:\Users\mdolf\OneDrive\Documents\Didattica
2026\London\Job\GitHub Materials\HousePrices.xlsx");

% Display results
HousePrices2
```

HousePrices2 = 47x3 table

	Sfeet	Bedrooms	Price
1	2104	3	399900

	Sfeet	Bedrooms	Price
2	1600	3	329900
3	2400	3	369000
4	1416	2	232000
5	3000	4	539900
6	1985	4	299900
7	1534	3	314900
8	1427	3	198999
9	1380	3	212000
10	1494	3	242500
11	1940	4	239999
12	2000	3	347000
13	1890	3	329999
14	4478	5	699900

⋮

Fit an OLS linear model using the table; by default the **last column** is treated as the response (set 'ResponseVar', 'Price' to be explicit).

```
model=fitlm(HousePrices)
```

```
model =  
Linear regression model:  
Price ~ 1 + Sfeet + Bedrooms
```


Estimated Coefficients:

	<u>Estimate</u>	<u>SE</u>	<u>tStat</u>	<u>pValue</u>
<b>(Intercept)</b>	1.0293e+05	52587	1.9573	0.056669
<b>Sfeet</b>	110.01	17.521	6.2784	1.3143e-07
<b>Bedrooms</b>	6948.8	19233	0.36129	0.71961

```
Number of observations: 47, Error degrees of freedom: 44  
Root Mean Squared Error: 8.33e+04  
R-squared: 0.576, Adjusted R-Squared: 0.556  
F-statistic vs. constant model: 29.9, p-value = 6.44e-09
```

Show the built-in diagnostic plot: with **one predictor it's data + fitted line**; with **multiple predictors it's an added-variable** (partial regression) plot.

```
plot(model)
```



Predict for a new case; when you pass a **matrix/row**, **columns must match the model's predictor order**. (Using a table with named variables avoids order mistakes.)

```
estimatedPrice= predict(model,[5000,1]);  
fprintf('Estimated price: $%.0f\n', estimatedPrice);
```

Estimated price: \$659908

## Theoretical background (example of unconstrained quadratic optimization).

Model and assumptions

Let  $y \in \mathbb{R}^n$  be the response (or regressand),  $X \in \mathbb{R}^{n \times p}$  the matrix of regressors,  $\beta \in \mathbb{R}^p$  the coefficients, and  $\varepsilon \in \mathbb{R}^n$  the errors.

$$y = X\beta + \varepsilon, \quad \mathbb{E}[\varepsilon] = 0, \quad \text{Var}(\varepsilon) = \sigma^2 I_n.$$

OLS minimizes the residual sum of squares (RSS):

$$\text{RSS}(\beta) = \|y - X\beta\|_2^2 = (y - X\beta)^\top (y - X\beta) = y^\top y - 2\beta^\top X^\top y + \beta^\top X^\top X \beta.$$

Since  $X^\top X \succcurlyeq 0$ ,  $\text{RSS}(\beta)$  is a convex quadratic in  $\beta$ . If  $\text{rank}(X) = p$ , the minimizer is unique.

Setting the gradient to zero yields the equations:

$$X^T X \hat{\beta} = X^T y \Rightarrow \hat{\beta} = (X^T X)^{-1} X^T y \quad (X^T X \text{ is invertible under Gauss-Markov assumptions}).$$

Under Gauss-Markov assumptions:

$$\hat{\sigma}^2 = \frac{\text{RSS}(\hat{\beta})}{n - p}, \quad \text{Var}(\hat{\beta}) = \sigma^2 (X^T X)^{-1} \approx \hat{\sigma}^2 (X^T X)^{-1}.$$

"Linear" refers to linearity in the coefficients, not necessarily in the raw predictors. A degree- $d$  polynomial regression augments predictors:

$$\phi(x) = [1, x, x^2, \dots, x^d],$$

and fits  $y = \phi(x)\beta + \varepsilon$  by OLS. The model is linear in  $\beta$  even though it is nonlinear in  $x$ . (In practice, center/scale predictors before adding higher-order terms to mitigate multicollinearity.)

RESET (linearity/specification)

To probe functional-form misspecification, augment the baseline model with powers of the fitted values, e.g.

$\hat{y}^2, \hat{y}^3$ , and test their joint significance. For nested models  $M_0$  (baseline) and  $M_1$  (augmented) with  $p_0 < p_1$

parameters and residual sums of squares  $\text{RSS}_0, \text{RSS}_1$ , the usual nested  $F$ -statistic is

$$F = \frac{(\text{RSS}_0 - \text{RSS}_1)/(p_1 - p_0)}{\text{RSS}_1/(n - p_1)}.$$

A small  $p$ -value suggests misspecification/nonlinearity in the baseline specification.