

QMF3_1_MachineLearning_Supervised

Logistic Regression for Binary Classification

This example illustrates how **logistic regression** can be used for **binary classification** problems.

Starting from raw exam score data, we:

- explore and visualize the dataset,
- train a logistic regression model using `fitglm`,
- interpret the model coefficients to understand the impact of each variable,
- compute predicted probabilities and training accuracy, and
- visualize the **decision boundary** separating the two classes.

The fitted model provides a simple yet powerful way to **estimate probabilities of categorical outcomes** and to **visualize the separating surface** in the predictor space — a key idea that extends naturally to higher-dimensional classification problems and more advanced machine learning models.

Data

The `data_classification.txt` contains scores for two exams in addition to a binary variable which denotes whether students were admitted to a university.

We obtain a logistic regression model to **predict admission probability** using the `fitglm` function.

Load the data into a table and preview the data

The first two columns (variables) will contain the exam scores and the third column the admission labels which we convert to `logical` values.

We also compute some summary statistics on the three variables.

```
clear;
data = readtable('data_classification.xlsx');

data.Properties.VariableNames = {'Exam1', 'Exam2', 'Admitted'};

data.Admitted = logical(data.Admitted)
```

data = 88×3 table

	Exam1	Exam2	Admitted
1	34	78	false
2	30	43	false
3	35	72	false
4	60	86	true
5	79	75	true
6	45	56	false
7	61	96	true
8	75	46	true
9	76	87	true
10	84	43	true
11	95	38	false
12	75	30	false
13	82	30	false
14	69	97	true

⋮

```
summary(data)
```

data: 88×3 table

Variables:

```
Exam1: double
Exam2: double
Admitted: logical (50 true)
```

Statistics for applicable variables:


	NumMissing	Min	Median	Max	Mean	Std
Exam1	0	30	67	99	65.2727	19.5093
Exam2	0	30	65.5000	98	64.2955	18.7673

We create a **grouped scatter plot** of the two exam scores, using the variable Admitted to define the groups.

- Each point represents a student.
- The x-axis shows **Exam 1** scores and the y-axis shows **Exam 2** scores.
- MATLAB automatically assigns different colors or markers for the two groups (Admitted = 1 and Admitted = 0).

This allows you to **visually separate admitted and non-admitted students** and observe whether a boundary between the two classes might exist in the exam-score space.

```
gscatter(data.Exam1,data.Exam2,data.Admitted)
```

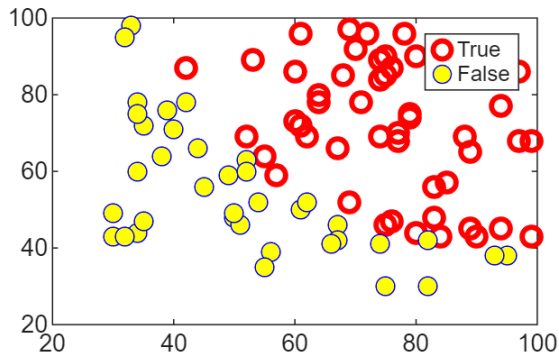


We now plot the data manually to have full control over the style and prepare for adding the decision boundary.

Unlike **gscatter**, this approach lets us customize marker colors, sizes, and legend entries — using red circles for admitted students and yellow-filled blue circles for those not admitted.

```
varNames = data.Properties.VariableNames;
% Plot the data with + for true and 0 for false examples
inds = data.(varNames{3}) == 1;
plot(data.(varNames{1})(inds), data.(varNames{2})(inds), 'ro', 'LineWidth', 2,
'MarkerSize', 7);
inds = data.(varNames{3}) == 0;
hold on

plot(data.(varNames{1})(inds), data.(varNames{2})(inds), 'bo', 'MarkerFaceColor',
'y', 'MarkerSize', 7);
legend('True', 'False')
hold off
```



Logistic Regression

We train the logistic regression model using the `fitglm` function, which fits a **generalized linear model** (GLM) to the data.

By setting the **'Distribution'** parameter to **'binomial'**, we specify that the response variable (Admitted) is **binary** — taking only the values 0 or 1.

This tells MATLAB to use the **logistic (logit) link function**, so the model estimates the **probability of admission** as a function of the two exam scores.

```
logMdl = fitglm(data, 'Distribution', 'binomial')
```

```
logMdl =  
Generalized linear regression model:  
  logit(Admitted) ~ 1 + Exam1 + Exam2  
  Distribution = Binomial
```

Estimated Coefficients:

	Estimate	SE	tStat	pValue
(Intercept)	-24.469	5.8101	-4.2115	2.5365e-05
Exam1	0.20664	0.049113	4.2074	2.5835e-05
Exam2	0.19073	0.04872	3.9147	9.0499e-05

```
88 observations, 85 error degrees of freedom  
Dispersion: 1  
Chi^2-statistic vs. constant model: 83.4, p-value = 7.89e-19
```

Note the form of the model displayed in the output above. This is short-hand for

$$\text{logit}(\text{Admitted}) = 1 * \theta_0 + \text{Exam1} * \theta_1 + \text{Exam2} * \theta_2$$

Since $\text{logit}(x)$ is the *inverse function* of $\text{sigmoid}(x)$, this model is equivalent to logistic regression model form for the probability of admission used in ex2:

$$\text{Admitted} = h_{\theta}(x) = \text{sigmoid}(\theta^T x),$$

where x includes the two exam scores and a bias term. A bias term is added automatically by `fitglm`.

Predict the training accuracy and probability of admission

Recall that a prediction of admission corresponds to a predicted probability > 0.5 . Run the code below to extract the θ values from the trained model, predict the probability of admission, and compute the training accuracy.

```
theta = logMdl.Coefficients.Estimate
```

```
theta = 3×1  
-24.4693  
 0.2066  
 0.1907
```

```
% Predict the probability for a student with scores of 45 and 85
```

```
prob = predict(logMdl,[45 85]);
```

```
fprintf('For a student with scores 45 and 85, we predict an admission probability  
of %f\n\n', prob);
```

```
For a student with scores 45 and 85, we predict an admission probability of 0.739104
```

```
% Compute the training accuracy
```

```
Admitted = predict(logMdl,data) > 0.5;
```

```
fprintf('Train Accuracy: %f\n', mean(double(Admitted == data.Admitted)) * 100);
```

```
Train Accuracy: 88.636364
```

Visualize the decision boundary

Create a grid of exam scores and create the decision boundary plot.

```
figure; hold on;
```

```
% Plot the positive and negative examples
```

```
varNames = data.Properties.VariableNames;
```

```
% Plot the data with + for true and 0 for false examples
```

```
inds = data.(varNames{3}) == 1;
```

```
plot(data.(varNames{1})(inds), data.(varNames{2})(inds), 'k+', 'LineWidth', 2,  
'MarkerSize', 7);
```

```
inds = data.(varNames{3}) == 0;
```

```
plot(data.(varNames{1})(inds), data.(varNames{2})(inds), 'ko', 'MarkerFaceColor',  
'y', 'MarkerSize', 7);
```

```
% Plot the decision boundary
```

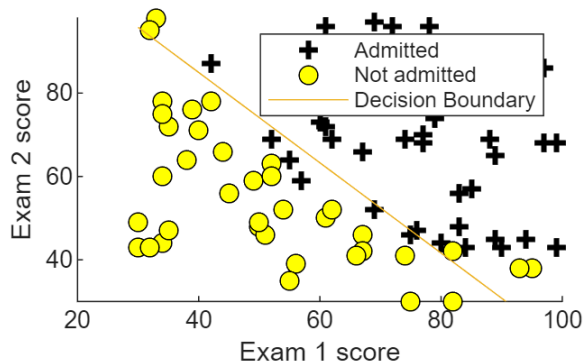
```
xvals = [min(data.Exam1), max(data.Exam1)];
```

```

yvals = -(theta(1)+theta(2)*xvals)/theta(3);
plot(xvals,yvals); hold off;
ylim([min(data.Exam2),max(data.Exam2)]);

% Labels and Legend
xlabel('Exam 1 score')
ylabel('Exam 2 score')
legend('Admitted','Not admitted','Decision Boundary')
hold off;

```



Next Step – Exercise on Classification (Exam Scores)

- Upload into the workspace the data contained into the spreadsheet data_classification.xlsx
- In the MATLAB Apps tab, select the Classification Learner app from the Machine Learning section
- Select 'New Session -> From Workspace' to start
- Under 'Workspace Variable', select 'data' and you will automatically find Exam1 and Exam2
- The app should correctly detect Admitted as the response variable to predict.
- Choose the default validation option.
- Select a model (e.g., logistic regression, decision tree, or SVM) and click the **Train** button.
- Experiment with several standard models using default settings and see if you can achieve at least 80% accuracy.

This exercise helps you compare the manual logistic regression workflow with MATLAB's interactive machine learning tools, reinforcing how supervised learning models can be trained and evaluated in different ways.